

Walk the pipeline, one gate at a time.

You don't need the tool. For staged teams, the 8-gate framework is a sequential checklist: at each gate ask the question, and a **"no"** routes you out to **real people** instead of letting an ungrounded synthetic run drive the decision. Answer all eight "yes" and you've earned a synthetic pass — still confirmed by humans at the ship point.

ASK AT EACH GATE — A "NO" IS WHERE YOU STOP

- 1 Do you have a clear, documented decision target?**
Vague or exploratory → **define the decision first**; scope it down to one answerable question.
- 2 Is it low-to-medium risk and reversible?**
Irreversible, or touches billing / contracts / security → **reframe** into smaller, reversible releases.
- 3 Is it a structured interaction or pricing-logic check?**
Speculative or generative discovery → **run human discovery first**; synthetics can't find net-new needs. [NN/g 2024](#)
- 4 Do you have live customer data to ground it?**
No telemetry / logs / interviews → **real customers only**; ungrounded synthetics hallucinate. [Making Science 2025](#)
- 5 Is the persona specific (role, rights, budget, expertise)?**
Generic template → **recalibrate**; there's no testable scope until the persona is concrete and bounded.
- 6 Does the agent replicate a known past result?**
No baseline, or drift > 15% → **recalibrate**; tighten prompts and re-test. [MeasuringU 2026](#)
- 7 Is a human + synthetic check-loop in place?**
Synthetic-only → **real customers only**; a closed loop never disagrees with you. [Codexical 2026](#)
- 8 Are compliance tags and drift reviews defined?**
No tagging → **brainstorm only**; untagged output isn't decision-grade. [Koji 2026](#)

Two gates are hard stops for humans

Gate 4 — no real data and **Gate 7 — no human loop** both route to **real customers only**. Grounding and human validation are the anchor; without either, a synthetic pass is just confident fiction.

Where a "no" routes you

- Define the decision / reframe (gates 1-2) — the question is too loose or too risky. Shrink it.
- Human discovery first (gate 3) — exploratory work needs live interviews, not simulation.
- Real customers only (gates 4 & 7) — no grounding data, or no human loop.
- Recalibrate (gates 5-6) — generic persona or failed baseline; fix and re-test.
- Brainstorm only (gate 8) — untagged output; use for divergent ideas, not validation.

Why the gates exist

Synthetic users are fast but systematically **over-agreeable, emotionally flat, and confidently wrong** on decision-grade work. The gates force grounding (gate 4), calibration (gate 6), and a human loop (gate 7) so a synthetic pass is earned, not assumed. The research consensus: **augment, never replace**. [Pehar 2025](#) · [Holli Downs 2025](#)