

Why we stop to bring in real people.

You don't need the tool to run Anchorflight — it's a decision discipline. At every turn of the build loop, ask one question: **can I trust a synthetic user here, or would its known biases corrupt this decision?** Synthetic users are fast and cheap, but the research is consistent: they are systematically over-agreeable, emotionally flat, and confidently wrong on decision-grade work. So you let them run where mistakes are cheap and reversible, and you stop for real humans where being wrong costs money, trust, or safety.

THE LOOP, AND WHAT EACH STAGE IS FOR

Frame

GO

Explore freely — synthetic is reliable.

Use synthetic users to generate hypotheses, learn a domain's vocabulary, and rehearse research guides. Just never let them **validate** the idea. [NN/g 2024](#) · [Koji 2026](#)

Generate

BUILD

Just build — no user claim is being made yet.

No trust decision here; you're producing the thing, not deciding it works.

Inspect

DECIDE

The pivot: is this reversible and low-stakes?

Yes → a synthetic gut-check is enough, keep moving. **No** (irreversible or high-stakes) → **stop**. Synthetic data is plausible but wrong on complex, decision-grade tasks. [MeasuringU 2026](#) · [Koji 2026](#)

Iterate

GO

Fast synthetic cycles for directional signals.

Fine for early, task-based usability and A/B exploration — not for discovery interviews or emotion-heavy work. [UXArmy 2025](#)

Ship

HUMANS

Validate with real people before you ship.

Anything touching pricing, onboarding, feature gates, personas, prioritization, or emotion needs human confirmation first. Synthetic participants never disagree with you. [Codexical 2026](#) · [Making Science 2025](#)

Start every decision with triage

Reversible & low-stakes → loose exit, trust the synthetic pass.

Irreversible & high-stakes → hard exit, real humans required. The stakes set the strictness before you run a single check.

Why we stop — the four failure modes

- ❑ **Over-agreeableness / sycophancy.** AI participants agree with your framing and rarely surface the objection that kills a feature. [Codexical 2026](#) · [Koji 2026](#)
- ❑ **Plausible but wrong.** Clean, confident output that doesn't match how real users behave on complex tasks. [MeasuringU 2026](#)
- ❑ **Emotional flatness & no context.** No real empathy or lived situation, so discovery and emotion-heavy work fall apart. [Making Science 2025](#) · [UXArmy 2025](#)
- ❑ **Structural bias.** Western defaults and assumption-echoing, so personas confirm instead of surprising you. [ACM Interactions](#) · [Pehar 2025](#)

Run it in 5 questions

- 1 Is this decision **reversible**?
- 2 Is it genuinely **low-stakes**?
- 3 Is it discovery, emotion, personas, or prioritization?
- 4 Would being **wrong** cost real money or trust?
- 5 Can real users **falsify** this before ship?

Any high-risk answer → stop the loop and bring in real people. The consensus across the research: augment, never replace. [Pehar 2025](#) · [Holli Downs 2025](#)